

Statistical Analysis Top 5 Leagues

November 19, 2024

1 Top 5 Leagues Historical Goals: Statistical Insight

1.1 Top 5 Leagues Historical Goals: France recent increase in goals

1.1.1 Import relevant libraries

```
[327]: import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns

from scipy import stats
from scipy.stats import linregress
from scipy.stats import norm
```

1.1.2 Load the CSV file with the data

```
[328]: # Read Top 5 Leagues 25 Countries CSV file
df = pd.read_csv(r"C:\Users\jprey\OneDrive\Escritorio\JP\KUL\5th Semester_\
↳(Polimi)\MIT IDSS\Practice Projects\Top 5_\
↳Leagues\top_5_leagues_25_countries_cumulative.csv")
df.head()
```

```
[328]: Countries 1963-1964 1964-1965 1965-1966 1966-1967 1967-1968 1968-1969 \
0 Germany 830 760 953 845 913 776
1 England 1132 1061 953 895 982 890
2 France 721 701 967 801 771 585
3 Spain 495 495 499 578 576 491
4 Italy 373 457 466 468 383 415

1969-1970 1970-1971 1971-1972 ... 2015-2016 2016-2017 2017-2018 \
0 881 869 940 ... 360 360 390
1 836 736 783 ... 297 291 275
2 699 759 788 ... 552 603 515
3 516 474 596 ... 573 647 584
4 396 423 427 ... 439 472 431
```

| | 2018-2019 | 2019-2020 | 2020-2021 | 2021-2022 | 2022-2023 | 2023-2024 | sum |
|---|-----------|-----------|-----------|-----------|-----------|-----------|-------|
| 0 | 384 | 364 | 351 | 365 | 419 | 440 | 39164 |
| 1 | 284 | 366 | 372 | 358 | 387 | 459 | 37995 |
| 2 | 614 | 493 | 614 | 700 | 681 | 568 | 37713 |
| 3 | 608 | 600 | 613 | 618 | 503 | 546 | 34409 |
| 4 | 417 | 424 | 374 | 392 | 310 | 321 | 27814 |

[5 rows x 63 columns]

```
[329]: # Create a sub-dataframe that includes only the 5 nations with the most overall
        ↪goals between 1994-2004
thirty_years_ago = df.iloc[:5, [0,32,33,34,35,36,37,38,39,40,41]]
# Add a column with its mean, use np.floor to avoid fractional goals.
thirty_years_ago['Mean'] = np.floor(thirty_years_ago.iloc[:, 1:].mean(axis=1))
# Make it an int.
thirty_years_ago["Mean"] = thirty_years_ago["Mean"].astype(int)

# Create a sub-dataframe that includes only the 5 nations with the most overall
        ↪goals between 2004-2014
previous_ten_years = df.iloc[:5, [0,42,43,44,45,46,47,48,49,50,51]]
# Add a column with its mean, use np.floor to avoid fractional goals.
previous_ten_years['Mean'] = np.floor(previous_ten_years.iloc[:, 1:].
        ↪mean(axis=1))
# Make it an int.
previous_ten_years["Mean"] = previous_ten_years["Mean"].astype(int)

# Create a sub-dataframe that includes only the 5 nations with the most overall
        ↪goals between 2014-2024
these_ten_years = df.iloc[:5, [0,52,53,54,55,56,57,58,59,60,61]]
# Add a column with its mean, use np.floor to avoid fractional goals.
these_ten_years['Mean'] = np.floor(these_ten_years.iloc[:, 1:].mean(axis=1))
# Make it an int.
these_ten_years["Mean"] = these_ten_years["Mean"].astype(int)
```

```
[330]: thirty_years_ago
```

```
[330]: Countries  1994-1995  1995-1996  1996-1997  1997-1998  1998-1999  1999-2000 \
0    Germany      695      593      572      563      508      445
1    England      779      604      541      489      437      433
2    France       670      616      671      568      565      558
3    Spain        558      734      638      402      469      559
4    Italy        487      544      568      611      489      463

        2000-2001  2001-2002  2002-2003  2003-2004  Mean
0          403      361      269      302      471
1          422      397      361      346      480
2          525      555      577      570      587
```

| | | | | | |
|---|-----|-----|-----|-----|-----|
| 3 | 617 | 560 | 564 | 552 | 565 |
| 4 | 503 | 518 | 538 | 518 | 523 |

```
[331]: previous_ten_years
```

```
[331]: Countries 2004-2005 2005-2006 2006-2007 2007-2008 2008-2009 2009-2010 \
0 Germany 317 327 302 292 295 312
1 England 367 355 323 327 310 342
2 France 482 438 443 519 503 457
3 Spain 513 491 457 500 587 574
4 Italy 635 626 673 664 616 569

2010-2011 2011-2012 2012-2013 2013-2014 Mean
0 348 365 437 390 338
1 320 323 290 320 327
2 467 507 584 531 493
3 513 546 626 619 542
4 522 475 488 474 574
```

```
[332]: these_ten_years
```

```
[332]: Countries 2014-2015 2015-2016 2016-2017 2017-2018 2018-2019 2019-2020 \
0 Germany 370 360 360 390 384 364
1 England 297 297 291 275 284 366
2 France 514 552 603 515 614 493
3 Spain 556 573 647 584 608 600
4 Italy 433 439 472 431 417 424

2020-2021 2021-2022 2022-2023 2023-2024 Mean
0 351 365 419 440 380
1 372 358 387 459 338
2 614 700 681 568 585
3 613 618 503 546 584
4 374 392 310 321 401
```

1.1.3 Extract the data relevant to France

```
[333]: # Extract the France values for the seasons 2004-2014 in a dataframe
france_previous_ten_years = previous_ten_years.iloc[2:3,
↪ [1,2,3,4,5,6,7,8,9,10]].iloc[0]
# Extract the France values for the seasons 2014-2024 in a dataframe
france_these_ten_years = these_ten_years.iloc[2:3, [1,2,3,4,5,6,7,8,9,10]].
↪ iloc[0]
```

1.1.4 T-test for two samples

We want to perform a t-test for two samples to check whether the mean goals of France have increased in these last 10 years compared to the previous last 10 years (2004-2014). To do so, we have to make sure that: - The samples are independent between each other and random - This is true because each year has different goals that do not depend on historical data. Moreover, the goals scored per season can be approximated to random since they depend on the players' performance year by year. - The samples are assumed to be normally distributed - A Shapiro-Wilk test will be performed to determine this. - The variances are equal. - Levene's test can be performed to determine this.

We will consider a level of significance of 0.01

```
[334]: # Set alpha as a global variable.
alpha = 0.01
```

Assess for normality: Shapiro-Wilk test For each sample: if the p-value is higher than alpha, then the null hypothesis stating that the sample follows a normal distribution cannot be rejected, i.e., the sample is normally distributed if $p\text{-value} > \alpha$.

```
[335]: # Shapiro test parameters for the previous last ten years (2004-2014)
shapiro_france_previous_ten_years = stats.shapiro(france_previous_ten_years)
p_value_previous_ten_years = shapiro_france_previous_ten_years.pvalue
print("Shapiro-Wilk Test for France's goals from 2004 to 2014: W =",
      ↪shapiro_france_previous_ten_years.statistic, ", p-value =",
      ↪p_value_previous_ten_years)

# Shapiro test parameters for this last ten years (2014-2024)
shapiro_france_these_ten_years = stats.shapiro(france_these_ten_years)
p_value_these_ten_years = shapiro_france_these_ten_years.pvalue
print("Shapiro-Wilk Test for France's goals from 2014 to 2024: W =",
      ↪shapiro_france_these_ten_years.statistic, ", p-value =",
      ↪p_value_these_ten_years)

# Conclusion
print("Both samples are normally distributed?", p_value_previous_ten_years >
      ↪alpha and p_value_these_ten_years > alpha)
```

```
Shapiro-Wilk Test for France's goals from 2004 to 2014: W = 0.9478259459792676 ,
p-value = 0.6428542654489622
```

```
Shapiro-Wilk Test for France's goals from 2014 to 2024: W = 0.9403216198196231 ,
p-value = 0.556597672066893
```

```
Both samples are normally distributed? True
```

Therefore, the samples are normally distributed.

Assess for equal variances: Levene's test We execute the test simultaneously for both samples: if the p-value is higher than alpha, then the null hypothesis stating that the variances of the

samples are equal cannot be rejected, i.e., the samples' variances are equal if $p\text{-value} > \alpha$.

```
[336]: # We execute Levene's test using our two datasets.
levene_test = stats.levene(france_previous_ten_years, france_these_ten_years)
p_value_levene_test = levene_test.pvalue
print("Levene's Test for both samples: Statistic =", levene_test.statistic, ", p-value =", p_value_levene_test)

# Conclusion
print("Both samples' variances are equal?", p_value_levene_test > alpha)
```

```
Levene's Test for both samples: Statistic = 2.358587346422444 , p-value = 0.14198606565081504
```

```
Both samples' variances are equal? True
```

Therefore, the samples' variances are equal.

The samples are independent, randomly selected, uniformly distributed and have equal variances.

1.1.5 Let's perform the t-test for two samples

- **Null Hypothesis (H₀):** previous \geq this (the mean of the goals scored by French players in the top 5 European Football Leagues between the seasons 2003-2024 and 2013-2024 is greater than or equal to the mean of the seasons 2013-2014 to 2023-2024)
- **Alternative Hypothesis (H_a):** previous $<$ this (the mean of the goals scored by French players in the top 5 European Football Leagues between the seasons 2003-2024 and 2013-2024 is less than the mean of the seasons 2013-2014 to 2023-2024)

```
[337]: # Using scipy module stats, a statistical test is performed.
t_statistic, p_value_France_Germany = stats.ttest_ind(france_previous_ten_years, france_these_ten_years, alternative="less")
print("p-value =", p_value_France_Germany)
print("Reject the null hypothesis in favor of the alternative hypothesis?", p_value_France_Germany < alpha)
```

```
p-value = 0.001274605152448159
```

```
Reject the null hypothesis in favor of the alternative hypothesis? True
```

1.1.6 We are Sure with a 99% Confidence Level that the mean of goals scored of French players in Europe Top 5 Football leagues has increased in these last 10 years with respect to the previous last 10 years (2004-2014)

1.2 Top 5 Leagues Historical Goals: Who will be at the top in the next decade

1.2.1 Let's interpret our findings regarding France.

France players have score more goals this last decade with respect to the previous one... does that mean that this trend will continue?

- Although computers and AI can do a lot for the Data Science domain, there is always the need for human insight. For example, it is true that France players have scored more this last decade, but does this mean that there will be a new increase the next decade? This is why it is very important to consider as much data as possible. In fact, this is why there is a different test for distributions of more than 30 samples.
- If we had only consider the previous result, we would have assumed an increase for the next decade, or we could have thought that it is reasonable to consider the average of the last decade only. However, this is a bad approximation. Goals scored by players per nation naturally vary throughout time.
- Modern football (the last 30 years) has become a more defensive game, leading to less goals. Moreover, globalization of the sport has brought great players from different countries, reducing the total sum of this historically frequent scoring countries. ##### This will be shown in the next graph

```
[338]: # Reshape the DataFrame using melt
df_melted = df.iloc[:5, :-1].melt(id_vars='Countries', var_name='Season',
    ↪value_name='Goals')

# Calculate the average for the first 30 years of the database and these last
    ↪30 years
average_60_to_30_years_ago_top_countries = df.iloc[0:5, 1:32].mean().mean()
average_30_to_present = df.iloc[0:5, 32:-1].mean().mean()
print("From 1963-1993, on average, the top 5 countries scored", int(np.
    ↪floor(average_60_to_30_years_ago_top_countries-average_30_to_present)),
    ↪"goals than in the past 30 years.")

# Create the line plot using Countries as hue
plt.figure(figsize=(12, 6))
sns.lineplot(data=df_melted, x='Season', y='Goals', hue='Countries', marker='o')

# Title and labels of the plot
plt.title('Goals per Season for the 5 Countries that Scored the Most in Europe
    ↪Top 5 Leagues (1994-2024)')
plt.xticks(rotation=-90)
plt.ylabel('Goals')
plt.xlabel('Season')


# Horizontal lines to show the average in their respective timeframe
plt.hlines(y=average_60_to_30_years_ago_top_countries, xmin="1963-1964",
    ↪xmax="1993-1994", label="Mean of goals from 60 to 30 years ago",
    ↪color='black', linestyle='-.', linewidth=4)
plt.hlines(y=average_30_to_present, xmin="1993-1994", xmax="2023-2024",
    ↪label="Mean of goals from 30 years ago until present", color='brown',
    ↪linestyle='-.', linewidth=4)

# Vertical dotted lines to show decade division
plt.axvline(x="1973-1974", color="black", linestyle='dotted')
```

```
plt.axvline(x="1983-1984", color="black", linestyle='dotted')
plt.axvline(x="1993-1994", color="black", linestyle='dotted')
plt.axvline(x="2003-2004", color="black", linestyle='dotted')
plt.axvline(x="2013-2014", color = "black", linestyle='dotted')

# Visualization
plt.grid()
plt.legend()
plt.tight_layout()
plt.show()
```

From 1963-1993, on average, the top 5 countries scored 198 goals than in the past 30 years.



As shown in the graph, there is a clear difference in these last 30 years compare to the previous ones.

- Moreover, the scoring trends in the last 30 years seems to have stabilized.

Which nation will be on top by the next decade?

- It would be interesting to consider the cumulative goals in the top 5 leagues since the birth of the most recent league (Bundesliga).

```
[339]: # Cumulative goals since 1963-1964. Table sorted by total sum
df.iloc[:, [0,62]].sort_values("sum", ascending = False).head()
```

```
[339]: Countries    sum
0    Germany  39164
```

| | | |
|---|---------|-------|
| 1 | England | 37995 |
| 2 | France | 37713 |
| 3 | Spain | 34409 |
| 4 | Italy | 27814 |

Germany is on top currently, but how likely is it for them to stay on the top for another decade? Let's dive deeper into the average goals scored per decade these last years.

```
[340]: # Concatenate all the means per decade to visualize it more easily.
means_table = pd.concat([thirty_years_ago.iloc[:, [0,-1]], previous_ten_years.
    ↳iloc[:, [-1]], these_ten_years.iloc[:, [-1]]], axis=1)
# Create an overall mean as int
means_table["Mean Modern Football"] = np.floor(means_table.iloc[:, 1:].
    ↳mean(axis=1))
means_table["Mean Modern Football"] = means_table["Mean Modern Football"].
    ↳astype(int)
# Add the total goals sum per country
means_table = pd.concat([means_table, df.iloc[0:5, [-1]]], axis=1)
# Rename the columns
means_table.columns = ["Countries", "Mean 30-20 years ago", "Mean 20-10 years_
    ↳ago", "Mean 10 years ago to present", "Overall Mean Past 30 years", "Total_
    ↳goals sum"]
# Show the table
means_table.reset_index(drop=True, inplace=True)
means_table
```

```
[340]: Countries Mean 30-20 years ago Mean 20-10 years ago \
0 Germany 471 338
1 England 480 327
2 France 587 493
3 Spain 565 542
4 Italy 523 574

Mean 10 years ago to present Overall Mean Past 30 years Total goals sum
0 380 396 39164
1 338 381 37995
2 585 555 37713
3 584 563 34409
4 401 499 27814
```

Who could dethrone Germany in the next decade? Let's analyze the Overall Mean of goals per season in the past 30 years and the historical sum of goals. - England has had less goals on average this past 30 years per season than Germany. They cannot beat Germany. - Italy is more than 12000 goals away, they cannot dethrone Germany. - Spain and France have scored on average the same amount of goals per season (± 8 goals) in these last 30 years. Therefore, the closest one to Germany is the best candidate. This is France

1.2.2 France vs Germany: race for gold

Let's perform a Normal Distribution problem where we provide the population standard deviations of means for each country. We will do a Linear Combination of random independent variables: -
 $\text{Var}(1F - 1G) = 1^2 \text{Var}(F) + (-1)^2 \text{Var}(G) = 1 \times \text{Var}(F) + 1 \times \text{Var}(G)$

We are evaluating over 30 different seasons, so z-tests can be performed

```
[341]: # Historical averages and population standard deviations. Also the current lead
        ↪ in accumulated goals is calculated
avg_france_goals_last_30 = means_table.iloc[2,4]
avg_germany_goals_last_30 = means_table.iloc[0,4]
std_france_goals_last_30 = df.iloc[2, 32:-1].std()
std_germany_goals_last_30 = df.iloc[4, 32:-1].std()
current_lead_germany = means_table.iloc[0,5] - means_table.iloc[2,5]

# Number of seasons to project
n_seasons = 10

# Linear combination of continuous random independent variables
mean_difference = avg_france_goals_last_30 - avg_germany_goals_last_30
std_france_and_germany = np.sqrt(std_france_goals_last_30**2 +
        ↪ std_germany_goals_last_30**2)

# Account for 10 years (it is NOT a linear combination, the variance gets
        ↪ multiplied by 10 and not 10 squared.)
mean_diff_10_years = mean_difference * n_seasons # Cumulative mean difference
        ↪ over 10 seasons
std_diff_10_years = std_france_and_germany * np.sqrt(n_seasons) # Cumulative
        ↪ standard deviation over 10 seasons

# Calculate the Z-score for the current lead of Germany
z_score = (mean_diff_10_years - current_lead_germany) / std_diff_10_years
probability = norm.cdf(z_score)
# Calculate the probability using the standard normal distribution
p_value_france_germany = round(1 - probability,4)
print(p_value_france_germany)

# Output the result
print(f"Probability that France will surpass Germany in {n_seasons} seasons:
        ↪ {probability*100}%")
print(f"Reject null hypothesis (France surpasses Germany)?",
        ↪ p_value_france_germany < alpha )
```

0.3538

Probability that France will surpass Germany in 10 seasons: 64.62122179514941%

Reject null hypothesis (France surpasses Germany)? False

We cannot conclude that France will surpass Germany with a 95% confidence level. The probability that France does surpass Germany is 64.62% ### Let's visualize the findings

```
[342]: # Create an array of x values from -4 to 4
x = np.linspace(-4, 4, 1000)
# Calculate the y values for the normal distribution
y = norm.pdf(x)

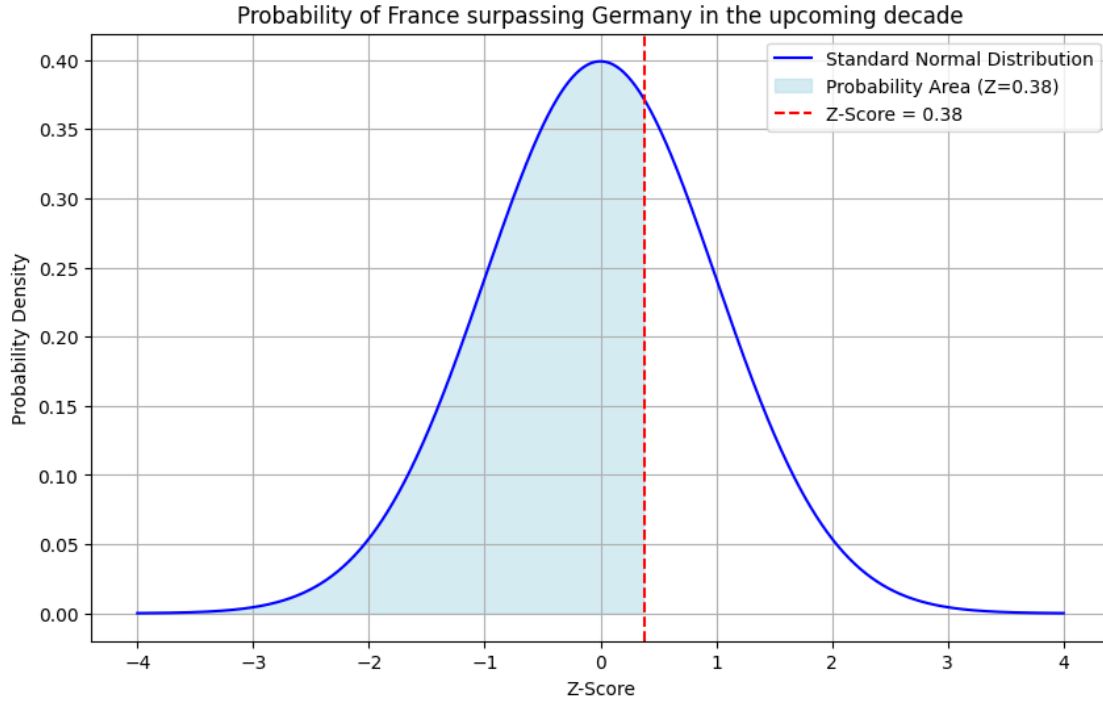
# Create the plot
plt.figure(figsize=(10, 6))
plt.plot(x, y, label='Standard Normal Distribution', color='blue')

# Shade the area under the curve to the right of the Z-score
plt.fill_between(x, y, where=(x <= z_score), color='lightblue', alpha=0.5,
    ↳label=f'Probability Area (Z={z_score:.2f})')

# Vertical line for the Z-score
plt.axvline(z_score, color='red', linestyle='--', label=f'Z-Score = {z_score:.
    ↳2f}')

# Add labels and title
plt.title('Probability of France surpassing Germany in the upcoming decade')
plt.xlabel('Z-Score')
plt.ylabel('Probability Density')
plt.legend()
plt.grid()

# Show the plot
plt.show()
```



1.2.3 Findings

- The whole point of this code is to show that it is easy to jump into wrong conclusions when not all the data is present.
- The increase of goals scored per season by France in the top 5 leagues in this decade with respect to the previous one is a fact (the null hypothesis was indeed rejected). However, this doesn't mean that the goals on next decade will also increase. Moreover, it doesn't even mean that the average will stay the same. It is necessary to account for as much data as possible.
- In this case, it was found that France has been rather monotonous this past 30 years in terms of goals scored per season in Europe's top 5 Football leagues.
- Furthermore, we were interested in predicting when is Germany going to be dethroned by most likely France. It was found that this is rather likely (64.62%) by the next decade. However, if the statistics follow the assumptions made, France will surpass Germany in 13 years and 7 months.